

Phase 2

Data preprocessing and analysis

Sama Amr -- 900211296 & Farida Madkour -- 900211360

```
In [1]: import warnings
warnings.filterwarnings('ignore')
###

%matplotlib inline
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [2]: df = pd.read_csv("bankloan.csv")
df.head()
```

```
Out[2]:
```

	ID	Age	Experience	Income	ZIP.Code	Family	CAvg	Education	Mortgage	Personal.Loan	!
0	1	25	1	49	91107	4	1.6	1	0	0	
1	2	45	19	34	90089	3	1.5	1	0	0	
2	3	39	15	11	94720	1	1.0	1	0	0	
3	4	35	9	100	94112	1	2.7	2	0	0	
4	5	35	8	45	91330	4	1.0	2	0	0	

```
In [3]: df.info()
print("-----")
print("List of Columns:", df.columns)
print("Shape:", df.shape)
```

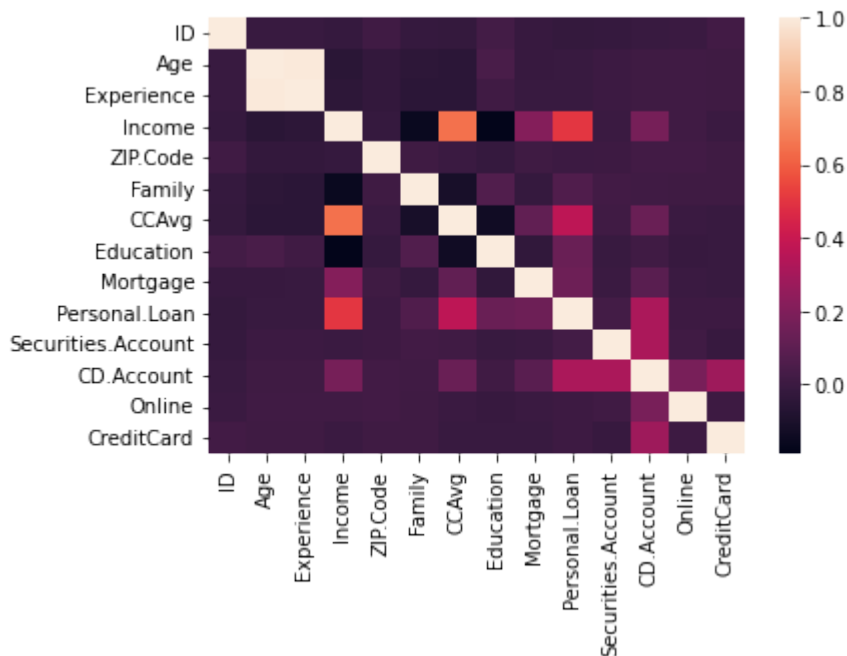
```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5000 entries, 0 to 4999
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---  ---
0   ID                    5000 non-null   int64
1   Age                   5000 non-null   int64
2   Experience             5000 non-null   int64
3   Income                5000 non-null   int64
4   ZIP.Code              5000 non-null   int64
5   Family                5000 non-null   int64
6   CCAvg                 5000 non-null   float64
7   Education             5000 non-null   int64
8   Mortgage              5000 non-null   int64
9   Personal.Loan         5000 non-null   int64
10  Securities.Account     5000 non-null   int64
11  CD.Account            5000 non-null   int64
12  Online                5000 non-null   int64
13  CreditCard            5000 non-null   int64
dtypes: float64(1), int64(13)
memory usage: 547.0 KB
-----
List of Columns: Index(['ID', 'Age', 'Experience', 'Income', 'ZIP.Code', 'Family',
                        'CCAvg', 'Education', 'Mortgage', 'Personal.Loan', 'Securities.Account',
                        'CD.Account', 'Online', 'CreditCard'],
                      dtype='object')
Shape: (5000, 14)

```

```
In [4]: sns.heatmap(df.corr(), annot=False)
```

```
Out[4]: <AxesSubplot:>
```



Drop ID, experience, and Zip Code columns since they're irrelevant

```
In [5]: df = df.drop(columns=['ID', 'Experience', 'ZIP.Code'])
df.head()
```

```
Out[5]:
```

	Age	Income	Family	CCAvg	Education	Mortgage	Personal.Loan	Securities.Account	CD.Acco
0	25	49	4	1.6	1	0	0		1
1	45	34	3	1.5	1	0	0		1
2	39	11	1	1.0	1	0	0		0
3	35	100	1	2.7	2	0	0		0
4	35	45	4	1.0	2	0	0		0

Check for missing values

```
In [6]: df.isnull().sum()
```

```
Out[6]:
```

Age	0
Income	0
Family	0
CCAvg	0
Education	0
Mortgage	0
Personal.Loan	0
Securities.Account	0
CD.Account	0
Online	0
CreditCard	0
dtype:	int64

Therefore, there is no missing values as specified by the non-null count and the sum calculated

Check for duplicate values and drop them

```
In [7]: df.duplicated().sum()
```

```
Out[7]: 13
```

```
In [8]: df.drop_duplicates(inplace=True)
df.duplicated().sum()
```

```
Out[8]: 0
```

Encodings

Change numeric/continous variables to type float and categorical/discrete variable to type category

```
In [9]: df['Income']=df['Income'].astype('float')
df['Family']=df['Family'].astype('category')
df['Education']=df['Education'].astype('category')
df['CCAvg']=df['CCAvg'].astype('float')
df['Mortgage']=df['Mortgage'].astype('float')
df['Personal.Loan']=df['Personal.Loan'].astype('category')
df['Securities.Account']=df['Securities.Account'].astype('category')
df['CD.Account']=df['CD.Account'].astype('category')
df['Online']=df['Online'].astype('category')
```

```
df['CreditCard']=df['CreditCard'].astype('category')
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 4987 entries, 0 to 4999
Data columns (total 11 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Age                    4987 non-null   int64
1   Income                 4987 non-null   float64
2   Family                 4987 non-null   category
3   CCAvg                  4987 non-null   float64
4   Education              4987 non-null   category
5   Mortgage               4987 non-null   float64
6   Personal.Loan          4987 non-null   category
7   Securities.Account      4987 non-null   category
8   CD.Account             4987 non-null   category
9   Online                 4987 non-null   category
10  CreditCard             4987 non-null   category
dtypes: category(7), float64(3), int64(1)
memory usage: 229.8 KB
```

Cut the Age and income into Ranges for better interpretations

```
In [10]: #minimum age = 23
#maximum age = 67
bins = [23,30,40,50,60,70]
df['Age_r'] = pd.cut(df['Age'], bins=bins, labels=['23-30', '30-40', '40-50', '50-60', '60-70'])

#minimum age = 8
#maximum age = 224
bins = [8,20,100,150,200,250]
df['Income_r'] = pd.cut(df['Income'], bins=bins, labels=['Poor', 'Middle_Class', 'Upper_Middle_Class', 'Very_High_Income'])
df.head()
```

```
Out[10]:
```

	Age	Income	Family	CCAvg	Education	Mortgage	Personal.Loan	Securities.Account	CD.Account
0	25	49.0	4	1.6	1	0.0	0	1	
1	45	34.0	3	1.5	1	0.0	0	1	
2	39	11.0	1	1.0	1	0.0	0	0	
3	35	100.0	1	2.7	2	0.0	0	0	
4	35	45.0	4	1.0	2	0.0	0	0	

Unique values of each of the variables

```
In [11]: print("Unique Family",pd.unique(df['Family']))
print("-----")
print("Unique Education",pd.unique(df['Education']))
print("-----")
print("Unique Personal.Loan",pd.unique(df['Personal.Loan']))
print("-----")
print("Unique Securities.Account",pd.unique(df['Securities.Account']))
print("-----")
print("Unique CD.Account",pd.unique(df['CD.Account']))
print("-----")
print("Unique Online",pd.unique(df['Online']))
print("-----")
print("Unique CreditCard",pd.unique(df['CreditCard']))
```

```
Unique Family [4, 3, 1, 2]
Categories (4, int64): [1, 2, 3, 4]
```

```
Unique Education [1, 2, 3]
Categories (3, int64): [1, 2, 3]
```

```
Unique Personal.Loan [0, 1]
Categories (2, int64): [0, 1]
```

```
Unique Securities.Account [1, 0]
Categories (2, int64): [0, 1]
```

```
Unique CD.Account [0, 1]
Categories (2, int64): [0, 1]
```

```
Unique Online [0, 1]
Categories (2, int64): [0, 1]
```

```
Unique CreditCard [0, 1]
Categories (2, int64): [0, 1]
```

Correlation Matrix along with a heatmap for our Numerical/Continuous Variables


```
In [12]: df.corr()
```

```
Out[12]:
```

	Age	Income	CCAvg	Mortgage
Age	1.000000	-0.056897	-0.052522	-0.013014
Income	-0.056897	1.000000	0.646065	0.206420
CCAvg	-0.052522	0.646065	1.000000	0.109162
Mortgage	-0.013014	0.206420	0.109162	1.000000

```
In [13]: sns.heatmap(df.corr(), annot=True)
```

```
Out[13]: <AxesSubplot:>
```



Barplots for our Discrete variables to show their distributions

```
In [14]: plt.figure(figsize=(30,28))
plt.subplot(3,4,1)
df['Family'].value_counts().plot(kind='bar')
plt.title("Family", fontsize=20, fontweight="bold")
plt.xlabel('Family', fontsize=15)
plt.ylabel('Count', fontsize=15)
```

```

plt.xticks(fontsize=15)
plt.yticks(fontsize=15)
counts=df['Family'].value_counts()
for i, count in enumerate (counts):
    plt.text(i,count+1,str(count), ha='center',va='bottom',fontsize=15)

plt.subplot(3,4,2)
df['Education'].value_counts().plot(kind='bar')
plt.title("Education", fontsize=20, fontweight="bold")
plt.xlabel('Education',fontsize=15)
plt.ylabel('Count',fontsize=15)
plt.xticks(fontsize=15)
plt.yticks(fontsize=15)
counts=df['Education'].value_counts()
for i, count in enumerate (counts):
    plt.text(i,count+1,str(count), ha='center',va='bottom',fontsize=15)

plt.subplot(3,4,3)
df['Personal.Loan'].value_counts().plot(kind='bar')
plt.title("Personal.Loan", fontsize=20, fontweight="bold")
plt.xlabel('Personal.Loan',fontsize=15)
plt.ylabel('Count',fontsize=15)
plt.xticks(fontsize=15)
plt.yticks(fontsize=15)
counts=df['Personal.Loan'].value_counts()
for i, count in enumerate (counts):
    plt.text(i,count+1,str(count), ha='center',va='bottom',fontsize=15)

plt.subplot(3,4,4)
df['Securities.Account'].value_counts().plot(kind='bar')
plt.title("Securities.Account", fontsize=20, fontweight="bold")
plt.xlabel('Securities.Account',fontsize=15)
plt.ylabel('Count',fontsize=15)
plt.xticks(fontsize=15)
plt.yticks(fontsize=15)
counts=df['Securities.Account'].value_counts()
for i, count in enumerate (counts):
    plt.text(i,count+1,str(count), ha='center',va='bottom',fontsize=15)

plt.subplot(3,4,5)
df['CD.Account'].value_counts().plot(kind='bar')
plt.title("CD.Account", fontsize=20, fontweight="bold")
plt.xlabel('CD.Account',fontsize=15)
plt.ylabel('Count',fontsize=15)
plt.xticks(fontsize=15)
plt.yticks(fontsize=15)
counts=df['CD.Account'].value_counts()
for i, count in enumerate (counts):
    plt.text(i,count+1,str(count), ha='center',va='bottom',fontsize=15)

plt.subplot(3,4,6)
df['Online'].value_counts().plot(kind='bar')
plt.title("Online", fontsize=20, fontweight="bold")
plt.xlabel('Online',fontsize=15)
plt.ylabel('Count',fontsize=15)
plt.xticks(fontsize=15)
plt.yticks(fontsize=15)
counts=df['Online'].value_counts()
for i, count in enumerate (counts):
    plt.text(i,count+1,str(count), ha='center',va='bottom',fontsize=15)

plt.subplot(3,4,7)
df['CreditCard'].value_counts().plot(kind='bar')
plt.title("CreditCard", fontsize=20, fontweight="bold")


```

```

plt.xlabel('CreditCard',fontSize=15)
plt.ylabel('Count',fontSize=15)
plt.xticks(fontsize=15)
plt.yticks(fontsize=15)
counts=df['CreditCard'].value_counts()
for i, count in enumerate (counts):
    plt.text(i,count+1,str(count), ha='center',va='bottom',fontSize=15)

plt.subplots_adjust(wspace=0.5,hspace=0.5)
plt.show()

```



Linegraphs for our Continous variables to show their distributions

```

In [15]: plt.figure(figsize=(30,15))
plt.subplot(2,2,1)
df['Age'].plot(kind='density')
plt.title("Age", fontsize=20, fontweight="bold")
plt.xlabel('Age',fontSize=20)
plt.ylabel('Density',fontSize=20)
plt.xticks(fontsize=15)
plt.yticks(fontsize=15)

plt.subplot(2,2,2)
df['Income'].plot(kind='density')
plt.xlabel('Income',fontSize=20)
plt.ylabel('Density',fontSize=20)
plt.title("Income", fontsize=20, fontweight="bold")
plt.xticks(fontsize=15)
plt.yticks(fontsize=15)

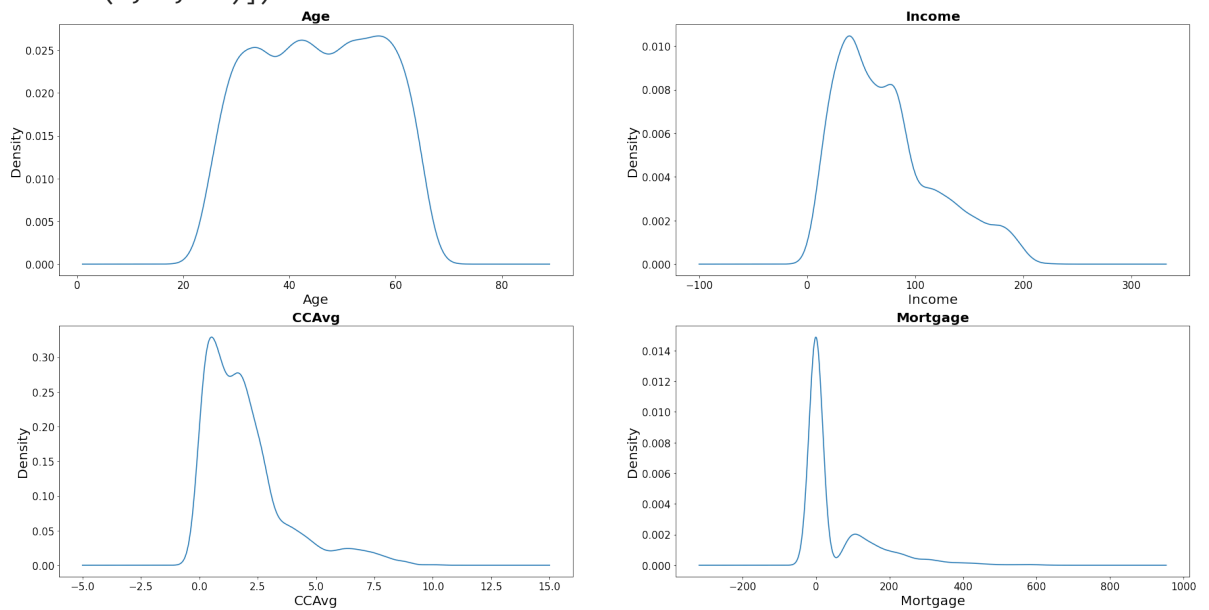
plt.subplot(2,2,3)
df['CCAvg'].plot(kind='density')
plt.title("CCAvg", fontsize=20, fontweight="bold")
plt.xlabel('CCAvg',fontSize=20)
plt.ylabel('Density',fontSize=20)
plt.xticks(fontsize=15)
plt.yticks(fontsize=15)

plt.subplot(2,2,4)
df['Mortgage'].plot(kind='density')
plt.title("Mortgage", fontsize=20, fontweight="bold")

```

```
plt.xlabel('Mortgage',fontsize=20)
plt.ylabel('Density',fontsize=20)
plt.xticks(fontsize=15)
plt.yticks(fontsize=15)
```

```
Out[15]: (array([-0.002, 0.    , 0.002, 0.004, 0.006, 0.008, 0.01 , 0.012,
        0.014, 0.016]),
 [Text(0, 0, ''),
  Text(0, 0, ''),
  Text(0, 0, ''),
  Text(0, 0, ''),
  Text(0, 0, ''),
  Text(0, 0, ''),
  Text(0, 0, ''),
  Text(0, 0, ''),
  Text(0, 0, ''),
  Text(0, 0, '')])
```



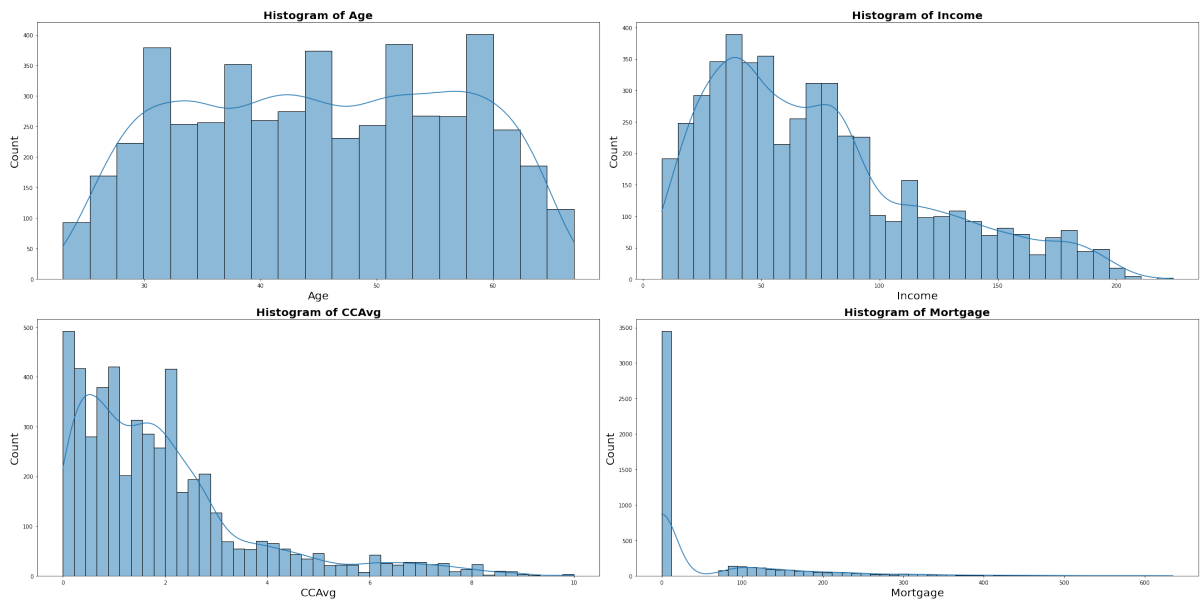
```
In [16]: dist_columns = ['Age', 'Income', 'CCAvg','Mortgage']

fig, axes = plt.subplots(nrows=2, ncols=2, figsize=(30, 15))

axes = axes.flatten()

for i, col in enumerate(dist_columns):
    sns.histplot(df[col], kde=True, ax=axes[i])
    axes[i].set_title(f'Histogram of {col}',fontsize=20,fontweight='bold')
    axes[i].set_xlabel(col,fontsize=20)
    axes[i].set_ylabel('Count',fontsize=20)

plt.tight_layout()
plt.show()
```

Box plots to show the relations and errors between each pair of variables

```
In [17]: plt.figure(figsize=(30,28))
plt.subplot(3,3,1)
sns.boxplot(x='Personal.Loan', y='Income', data=df, palette='viridis')
plt.xlabel('Personal.Loan',fontsize=15)
plt.ylabel('Income',fontsize=15)
plt.title('Personal.Loan x Income',fontsize=15, fontweight='bold')
plt.xticks(fontsize=15)
plt.yticks(fontsize=15)

plt.subplot(3,3,2)
sns.boxplot(x='CreditCard', y='Income', data=df, palette='viridis')
plt.xlabel('CreditCard',fontsize=15)
plt.ylabel('Income',fontsize=15)
plt.title('CreditCard x Income',fontsize=15, fontweight='bold')
plt.xticks(fontsize=15)
plt.yticks(fontsize=15)

plt.subplot(3,3,3)
sns.boxplot(x='Income_r', y='Mortgage', data=df, palette='viridis')
plt.xlabel('Income_range',fontsize=15)
plt.ylabel('Mortgage',fontsize=15)
plt.title('Income_range x Mortgage',fontsize=15, fontweight='bold')
plt.xticks(fontsize=15)
plt.yticks(fontsize=15)

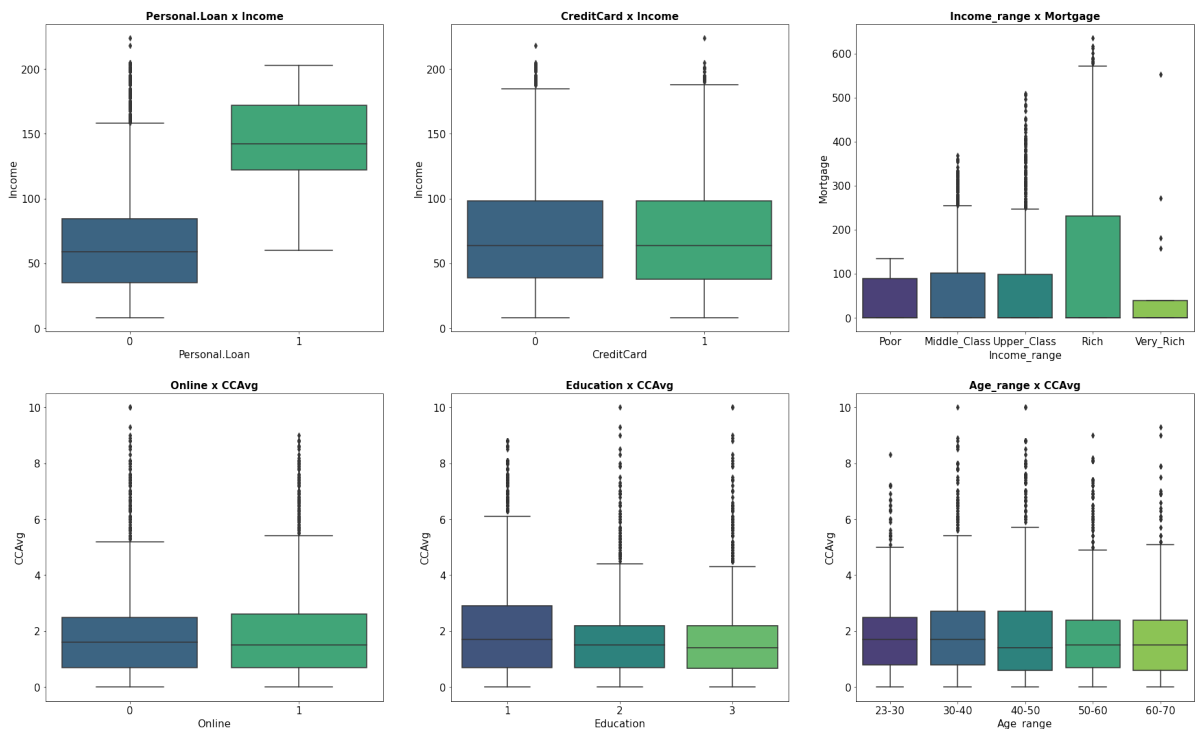
plt.subplot(3,3,4)
sns.boxplot(x='Online', y='CCAvg', data=df, palette='viridis')
plt.xlabel('Online',fontsize=15)
plt.ylabel('CCAvg',fontsize=15)
plt.title('Online x CCAvg',fontsize=15, fontweight='bold')
plt.xticks(fontsize=15)
plt.yticks(fontsize=15)

plt.subplot(3,3,5)
sns.boxplot(x='Education', y='CCAvg', data=df, palette='viridis')
plt.xlabel('Education',fontsize=15)
plt.ylabel('CCAvg',fontsize=15)
plt.title('Education x CCAvg',fontsize=15, fontweight='bold')
plt.xticks(fontsize=15)
plt.yticks(fontsize=15)

plt.subplot(3,3,6)
```

```
sns.boxplot(x='Age_r', y='CCAvg', data=df, palette='viridis')
plt.xlabel('Age_range',fontsize=15)
plt.ylabel('CCAvg',fontsize=15)
plt.title('Age_range x CCAvg',fontsize=15, fontweight='bold')
plt.xticks(fontsize=15)
plt.yticks(fontsize=15)

plt.show()
```



We can identify that income, CCAvg and Mortgage are 3 features containing outliers

Function to identify outliers

```
In [18]: def outlier(df):

    Q1=df.quantile(0.25)

    Q3=df.quantile(0.75)

    IQR=Q3-Q1

    out = df[((df<(Q1-1.5*IQR)) | (df>(Q3+1.5*IQR)))]

    return out
```

```
In [19]: skewed=['Income','CCAvg','Mortgage']
for col in skewed:
    outliers=outlier(df[col])
    print("Number of outliers in",col,":", str(len(outliers)),",It's Percentage is
    print("\n")
```

Number of outliers in Income : 96 ,It's Percentage is : 1.925005013033888 %

Number of outliers in CCAvg : 301 ,It's Percentage is : 6.035692801283337 %

Number of outliers in Mortgage : 291 ,It's Percentage is : 5.835171445758974 %

Outlier numbers are relatively low, yet they could better. In addition their line graphs and histograms are skewed. We found a solution to the problems by:

Find a suitable transformation for the skewed features

```
In [20]: skewed = ['Income', 'CCAvg', 'Mortgage']

for col in skewed:
    fig, axes = plt.subplots(1, 4, figsize=(20, 5))

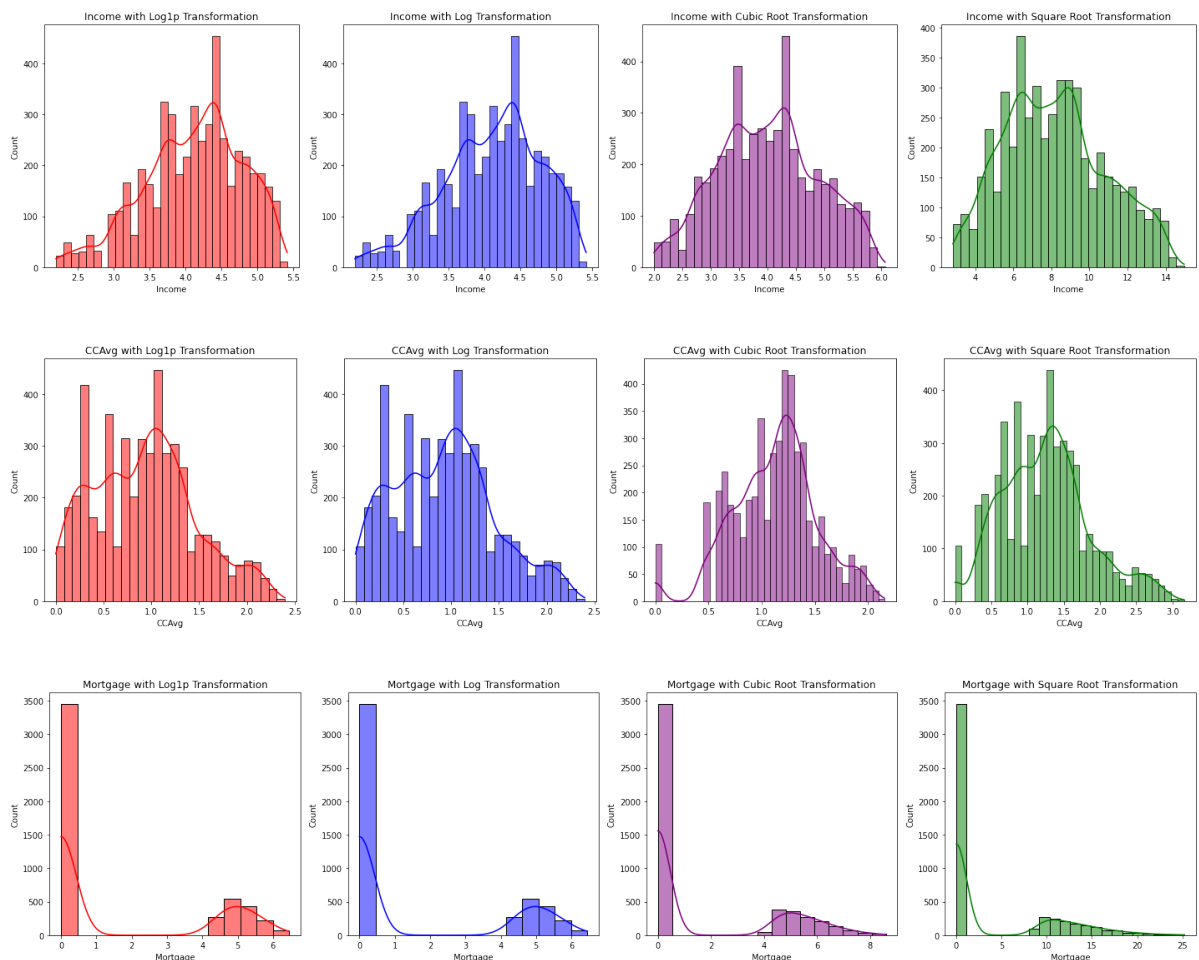
    # Log1p Transformation
    sns.histplot(np.log1p(df[col]), color='red', ax=axes[0], kde=True)
    axes[0].set_title(f'{col} with Log1p Transformation')

    # Log Transformation
    sns.histplot(np.log(df[col] + 1), color='blue', ax=axes[1], kde=True)
    axes[1].set_title(f'{col} with Log Transformation')

    # Cubic Root Transformation
    sns.histplot(np.cbrt(df[col]), color='purple', ax=axes[2], kde=True)
    axes[2].set_title(f'{col} with Cubic Root Transformation')

    # Square Root Transformation
    sns.histplot(np.sqrt(df[col]), color='green', ax=axes[3], kde=True)
    axes[3].set_title(f'{col} with Square Root Transformation')

plt.tight_layout()
plt.show()
print("\n")
```



Therefore, the best suitable transformation for:

- Income: Cubic root
- CCAvg: Cubic root
- Mortgage: Square root

```
In [21]: df['Income'] = np.cbrt(df['Income'])
df['CAvg'] = np.cbrt(df['CAvg'])
df['Mortgage'] = np.sqrt(df['Mortgage'])
```

Test for outliers after the transformation and plot the histograms

```
In [22]: skewed = ['Income', 'CAvg', 'Mortgage']
for col in skewed:
    outliers=outlier(df[col])
    print("Number of outliers in",col,":", str(len(outliers)),",It's Percentage is
    print("\n")
```

Number of outliers in Income : 0 ,It's Percentage is : 0.0 %

Number of outliers in CCAvg : 109 ,It's Percentage is : 2.1856827752155605 %

Number of outliers in Mortgage : 1 ,It's Percentage is : 0.020052135552436335 %

Outliers are significantly reduced after the transformation

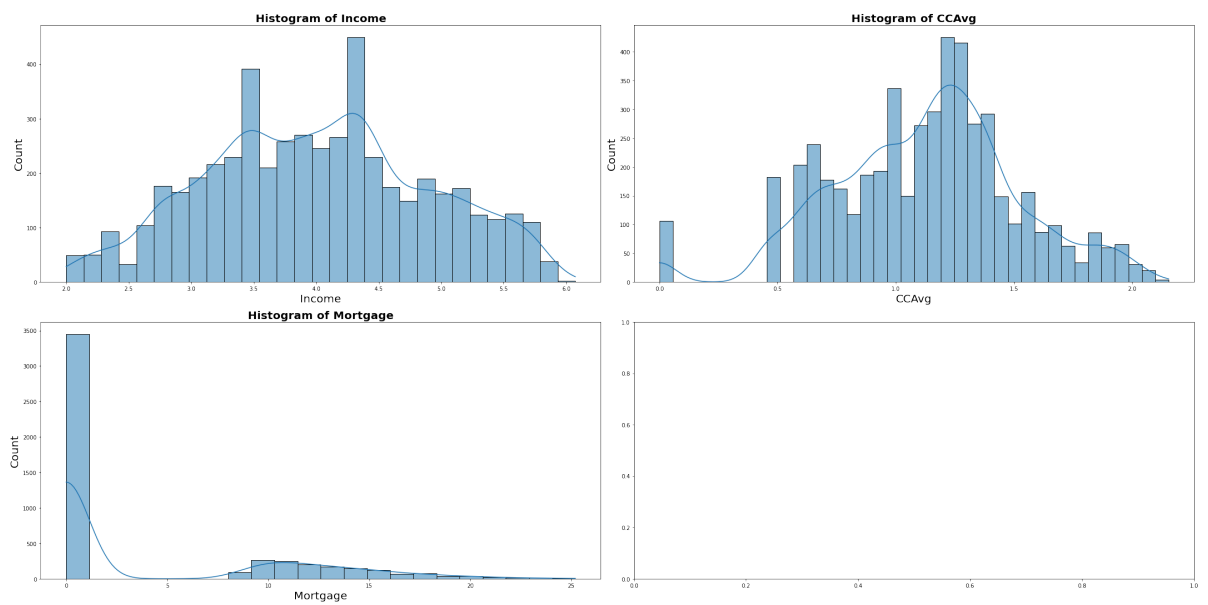
```
In [23]: dist_columns = ['Income', 'CAvg', 'Mortgage']

fig, axes = plt.subplots(nrows=2, ncols=2, figsize=(30, 15))

axes = axes.flatten()

for i, col in enumerate(dist_columns):
    sns.histplot(df[col], kde=True, ax=axes[i])
    axes[i].set_title(f'Histogram of {col}', fontsize=20, fontweight='bold')
    axes[i].set_xlabel(col, fontsize=20)
    axes[i].set_ylabel('Count', fontsize=20)

plt.tight_layout()
plt.show()
```



Finally, here's a summary of our continuous features

```
In [24]: df.describe()
```

```
Out[24]:
```

	Age	Income	CCAvg	Mortgage
count	4987.000000	4987.000000	4987.000000	4987.000000
mean	45.347704	4.006779	1.127270	4.046775
std	11.460838	0.887801	0.392599	6.346477
min	23.000000	2.000000	0.000000	0.000000
25%	35.000000	3.391211	0.887904	0.000000
50%	45.000000	4.000000	1.144714	0.000000
75%	55.000000	4.610436	1.375069	10.049876
max	67.000000	6.073178	2.154435	25.199206