

# Phase 3

Sama Amr -- 900211296 & Farida Madkour -- 900211360

```
In [ ]: import warnings
warnings.filterwarnings('ignore')
###

%matplotlib inline
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import cross_val_predict, KFold
from sklearn.neighbors import KNeighborsClassifier
from sklearn import metrics
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
```

```
In [ ]: from google.colab import files
uploaded = files.upload()
```


No file chosen      Upload widget is only available when the cell has been executed in the current browser session. Please rerun this cell to enable.  
Saving bankloan.csv to bankloan.csv

```
In [ ]: df = pd.read_csv("bankloan.csv")
df.head()
```

```
Out[ ]:
```

	ID	Age	Experience	Income	ZIP.Code	Family	CCAvg	Education	Mortgage	Personal.Loan	!
0	1	25	1	49	91107	4	1.6	1	0	0	
1	2	45	19	34	90089	3	1.5	1	0	0	
2	3	39	15	11	94720	1	1.0	1	0	0	
3	4	35	9	100	94112	1	2.7	2	0	0	
4	5	35	8	45	91330	4	1.0	2	0	0	

```
In [ ]: df.info()
print("-----")
print("List of Columns:", df.columns)
print("Shape:", df.shape)
```



Drop ID, experience, and Zip Code columns since they're irrelevant

```
In [ ]: df = df.drop(columns=['ID', 'Experience', 'ZIP.Code'])
df.head()
```

```
Out[ ]:
```

	Age	Income	Family	CCAvg	Education	Mortgage	Personal.Loan	Securities.Account	CD.Acco
0	25	49	4	1.6	1	0	0	1	
1	45	34	3	1.5	1	0	0	1	
2	39	11	1	1.0	1	0	0	0	
3	35	100	1	2.7	2	0	0	0	
4	35	45	4	1.0	2	0	0	0	

Check for missing values

```
In [ ]: df.isnull().sum()
```

```
Out[ ]:
```

Age	0
Income	0
Family	0
CCAvg	0
Education	0
Mortgage	0
Personal.Loan	0
Securities.Account	0
CD.Account	0
Online	0
CreditCard	0

dtype: int64

Therefore, there is no missing values as specified by the non-null count and the sum calculated

Check for duplicate values and drop them

```
In [ ]: df.duplicated().sum()
```

```
Out[ ]: 13
```

```
In [ ]: df.drop_duplicates(inplace=True)
df.duplicated().sum()
```

```
Out[ ]: 0
```

Encodings

Change numeric/continous variables to type float and categorical/discrete variable to type category

```
In [ ]: df['Income']=df['Income'].astype('float')
df['Family']=df['Family'].astype('category')
df['Education']=df['Education'].astype('category')
df['CCAvg']=df['CCAvg'].astype('float')
df['Mortgage']=df['Mortgage'].astype('float')
```

```
df['Personal.Loan']=df['Personal.Loan'].astype('category')
df['Securities.Account']=df['Securities.Account'].astype('category')
df['CD.Account']=df['CD.Account'].astype('category')
df['Online']=df['Online'].astype('category')
df['CreditCard']=df['CreditCard'].astype('category')
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 4987 entries, 0 to 4999
Data columns (total 11 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Age                   4987 non-null  int64
1   Income                4987 non-null  float64
2   Family                4987 non-null  category
3   CCAvg                 4987 non-null  float64
4   Education              4987 non-null  category
5   Mortgage              4987 non-null  float64
6   Personal.Loan         4987 non-null  category
7   Securities.Account     4987 non-null  category
8   CD.Account            4987 non-null  category
9   Online                4987 non-null  category
10  CreditCard            4987 non-null  category
dtypes: category(7), float64(3), int64(1)
memory usage: 229.8 KB
```

Cut the Age and income into Ranges for better interpretations

```
In [ ]: #minimum age = 23
#maximum age = 67
bins = [22,30,40,50,60,70]
df['Age_r'] = pd.cut(df['Age'], bins=bins, labels=['23-30', '30-40', '40-50', '50-60', '60-70'])

#minimum age = 8
#maximum age = 224
bins = [7,20,100,150,200,250]
df['Income_r'] = pd.cut(df['Income'], bins=bins, labels=['Poor', 'Middle_Class', 'Upper_Middle_Class', 'Very_High_Income'])
df.head()
```

```
Out[ ]:
```

	Age	Income	Family	CCAvg	Education	Mortgage	Personal.Loan	Securities.Account	CD.Account
0	25	49.0	4	1.6	1	0.0	0	1	
1	45	34.0	3	1.5	1	0.0	0	1	
2	39	11.0	1	1.0	1	0.0	0	0	
3	35	100.0	1	2.7	2	0.0	0	0	
4	35	45.0	4	1.0	2	0.0	0	0	

Unique values of each of the variables

```
In [ ]: print("Unique Family",pd.unique(df['Family']))
print("-----")
print("Unique Education",pd.unique(df['Education']))
print("-----")
print("Unique Personal.Loan",pd.unique(df['Personal.Loan']))
print("-----")
print("Unique Securities.Account",pd.unique(df['Securities.Account']))
print("-----")
print("Unique CD.Account",pd.unique(df['CD.Account']))
```

```
print("-----")
print("Unique Online",pd.unique(df['Online']))
print("-----")
print("Unique CreditCard",pd.unique(df['CreditCard']))
```

```
Unique Family [4, 3, 1, 2]
Categories (4, int64): [1, 2, 3, 4]
```

```
-----
Unique Education [1, 2, 3]
Categories (3, int64): [1, 2, 3]
```

```
-----
Unique Personal.Loan [0, 1]
Categories (2, int64): [0, 1]
```

```
-----
Unique Securities.Account [1, 0]
Categories (2, int64): [0, 1]
```

```
-----
Unique CD.Account [0, 1]
Categories (2, int64): [0, 1]
```

```
-----
Unique Online [0, 1]
Categories (2, int64): [0, 1]
```

```
-----
Unique CreditCard [0, 1]
Categories (2, int64): [0, 1]
```

Correlation Matrix along with a heatmap for our Numerical/Continuous Variables

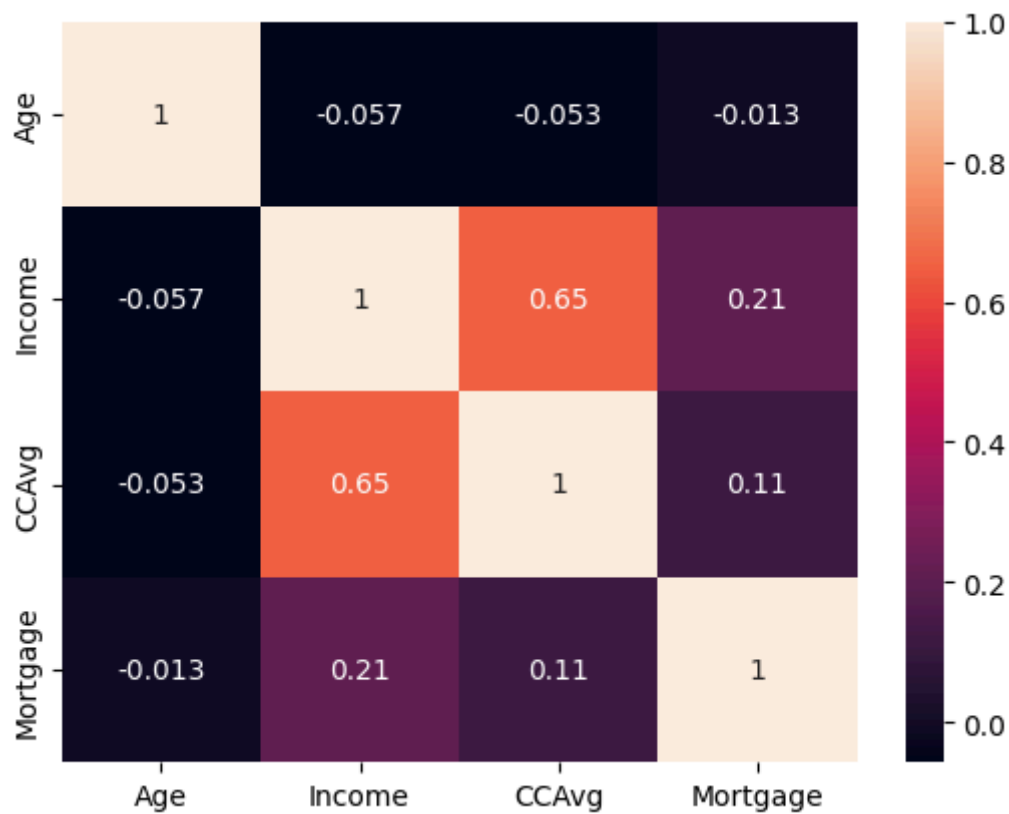
```
In [ ]: df.corr()
```

```
Out[ ]:
```

	Age	Income	CCAvg	Mortgage
Age	1.000000	-0.056897	-0.052522	-0.013014
Income	-0.056897	1.000000	0.646065	0.206420
CCAvg	-0.052522	0.646065	1.000000	0.109162
Mortgage	-0.013014	0.206420	0.109162	1.000000

```
In [ ]: sns.heatmap(df.corr(), annot=True)
```

```
Out[ ]: <Axes: >
```



Barplots for our Discrete variables to show their distributions

```
In [ ]: plt.figure(figsize=(30,28))
plt.subplot(3,4,1)
df['Family'].value_counts().plot(kind='bar')
plt.title("Family", fontsize=20, fontweight="bold")
plt.xlabel('Family',fontsize=15)
plt.ylabel('Count',fontsize=15)
plt.xticks(fontsize=15)
plt.yticks(fontsize=15)
counts=df['Family'].value_counts()
for i, count in enumerate (counts):
    plt.text(i,count+1,str(count), ha='center',va='bottom',fontsize=15)

plt.subplot(3,4,2)
df['Education'].value_counts().plot(kind='bar')
plt.title("Education", fontsize=20, fontweight="bold")
plt.xlabel('Education',fontsize=15)
plt.ylabel('Count',fontsize=15)
plt.xticks(fontsize=15)
plt.yticks(fontsize=15)
counts=df['Education'].value_counts()
for i, count in enumerate (counts):
    plt.text(i,count+1,str(count), ha='center',va='bottom',fontsize=15)

plt.subplot(3,4,3)
df['Personal.Loan'].value_counts().plot(kind='bar')
plt.title("Personal.Loan", fontsize=20, fontweight="bold")
plt.xlabel('Personal.Loan',fontsize=15)
plt.ylabel('Count',fontsize=15)
plt.xticks(fontsize=15)
plt.yticks(fontsize=15)
counts=df['Personal.Loan'].value_counts()
for i, count in enumerate (counts):
    plt.text(i,count+1,str(count), ha='center',va='bottom',fontsize=15)

plt.subplot(3,4,4)
```

```

df['Securities.Account'].value_counts().plot(kind='bar')
plt.title("Securities.Account", fontsize=20, fontweight="bold")
plt.xlabel('Securities.Account', fontsize=15)
plt.ylabel('Count', fontsize=15)
plt.xticks(fontsize=15)
plt.yticks(fontsize=15)
counts=df['Securities.Account'].value_counts()
for i, count in enumerate (counts):
    plt.text(i, count+1, str(count), ha='center', va='bottom', fontsize=15)


plt.subplot(3,4,5)
df['CD.Account'].value_counts().plot(kind='bar')
plt.title("CD.Account", fontsize=20, fontweight="bold")
plt.xlabel('CD.Account', fontsize=15)
plt.ylabel('Count', fontsize=15)
plt.xticks(fontsize=15)
plt.yticks(fontsize=15)
counts=df['CD.Account'].value_counts()
for i, count in enumerate (counts):
    plt.text(i, count+1, str(count), ha='center', va='bottom', fontsize=15)

plt.subplot(3,4,6)
df['Online'].value_counts().plot(kind='bar')
plt.title("Online", fontsize=20, fontweight="bold")
plt.xlabel('Online', fontsize=15)
plt.ylabel('Count', fontsize=15)
plt.xticks(fontsize=15)
plt.yticks(fontsize=15)
counts=df['Online'].value_counts()
for i, count in enumerate (counts):
    plt.text(i, count+1, str(count), ha='center', va='bottom', fontsize=15)

plt.subplot(3,4,7)
df['CreditCard'].value_counts().plot(kind='bar')
plt.title("CreditCard", fontsize=20, fontweight="bold")
plt.xlabel('CreditCard', fontsize=15)
plt.ylabel('Count', fontsize=15)
plt.xticks(fontsize=15)
plt.yticks(fontsize=15)
counts=df['CreditCard'].value_counts()
for i, count in enumerate (counts):
    plt.text(i, count+1, str(count), ha='center', va='bottom', fontsize=15)

plt.subplots_adjust(wspace=0.5, hspace=0.5)
plt.show()

```



Linegraphs for our Continuous variables to show their distributions

```
In [ ]: plt.figure(figsize=(30,15))
plt.subplot(2,2,1)
df['Age'].plot(kind='density')
plt.title("Age", fontsize=20, fontweight="bold")
plt.xlabel('Age', fontsize=20)
plt.ylabel('Density', fontsize=20)
plt.xticks(fontsize=15)
plt.yticks(fontsize=15)

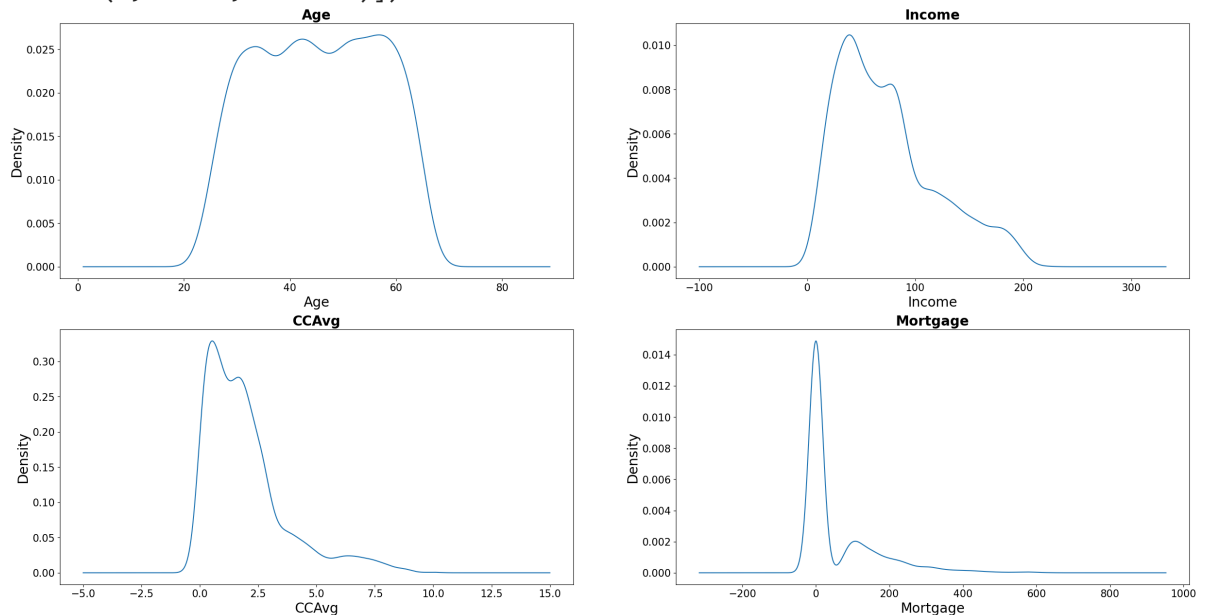
plt.subplot(2,2,2)
df['Income'].plot(kind='density')
plt.xlabel('Income', fontsize=20)
plt.ylabel('Density', fontsize=20)
plt.title("Income", fontsize=20, fontweight="bold")
plt.xticks(fontsize=15)
plt.yticks(fontsize=15)

plt.subplot(2,2,3)
df['CCAvg'].plot(kind='density')
plt.title("CCAvg", fontsize=20, fontweight="bold")
plt.xlabel('CCAvg', fontsize=20)
plt.ylabel('Density', fontsize=20)
plt.xticks(fontsize=15)
plt.yticks(fontsize=15)

plt.subplot(2,2,4)
df['Mortgage'].plot(kind='density')
plt.title("Mortgage", fontsize=20, fontweight="bold")
plt.xlabel('Mortgage', fontsize=20)
plt.ylabel('Density', fontsize=20)
plt.xticks(fontsize=15)
plt.yticks(fontsize=15)
```



```
Out[ ]: (array([-0.002, 0.    , 0.002, 0.004, 0.006, 0.008, 0.01 , 0.012,
                0.014, 0.016]),
        [Text(0, -0.002, '-0.002'),
         Text(0, 0.0, '0.000'),
         Text(0, 0.002, '0.002'),
         Text(0, 0.004, '0.004'),
         Text(0, 0.006, '0.006'),
         Text(0, 0.008, '0.008'),
         Text(0, 0.01, '0.010'),
         Text(0, 0.012, '0.012'),
         Text(0, 0.014, '0.014'),
         Text(0, 0.016, '0.016')])
```



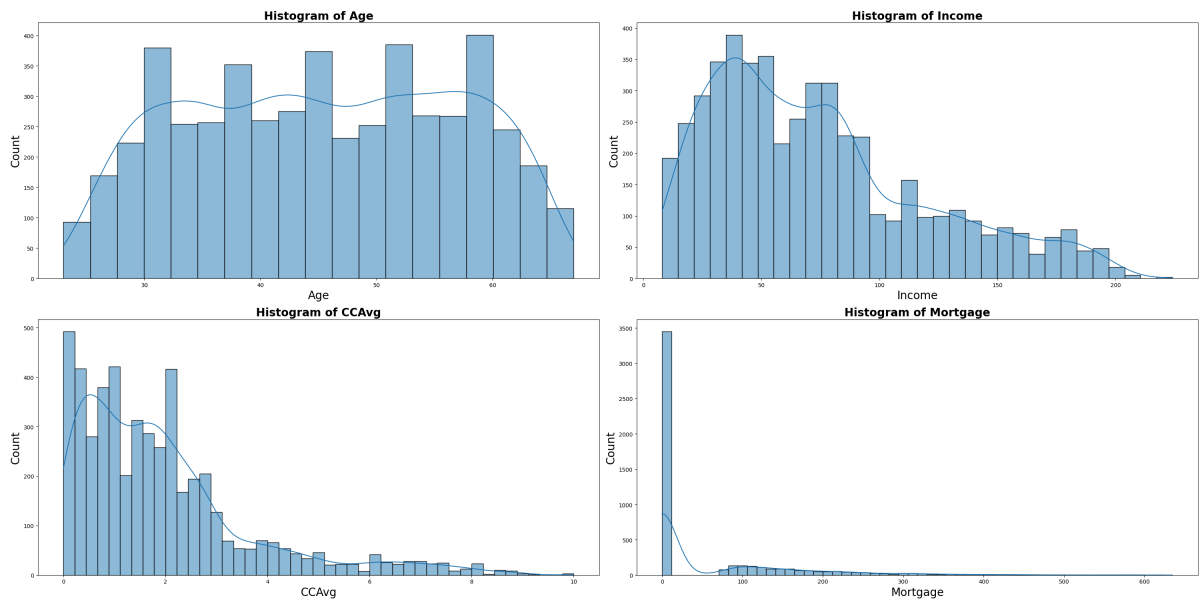
```
In [ ]: dist_columns = ['Age', 'Income', 'CCAvg', 'Mortgage']

fig, axes = plt.subplots(nrows=2, ncols=2, figsize=(30, 15))

axes = axes.flatten()

for i, col in enumerate(dist_columns):
    sns.histplot(df[col], kde=True, ax=axes[i])
    axes[i].set_title(f'Histogram of {col}', fontsize=20, fontweight='bold')
    axes[i].set_xlabel(col, fontsize=20)
    axes[i].set_ylabel('Count', fontsize=20)

plt.tight_layout()
plt.show()
```



Box plots to show the relations and errors between each pair of variables

```
In [ ]: plt.figure(figsize=(30,28))
plt.subplot(3,3,1)
sns.boxplot(x='Personal.Loan', y='Income', data=df, palette='viridis')
plt.xlabel('Personal.Loan',fontsize=15)
plt.ylabel('Income',fontsize=15)
plt.title('Personal.Loan x Income',fontsize=15, fontweight='bold')
plt.xticks(fontsize=15)
plt.yticks(fontsize=15)

plt.subplot(3,3,2)
sns.boxplot(x='CreditCard', y='Income', data=df, palette='viridis')
plt.xlabel('CreditCard',fontsize=15)
plt.ylabel('Income',fontsize=15)
plt.title('CreditCard x Income',fontsize=15, fontweight='bold')
plt.xticks(fontsize=15)
plt.yticks(fontsize=15)

plt.subplot(3,3,3)
sns.boxplot(x='Income_r', y='Mortgage', data=df, palette='viridis')
plt.xlabel('Income_range',fontsize=15)
plt.ylabel('Mortgage',fontsize=15)
plt.title('Income_range x Mortgage',fontsize=15, fontweight='bold')
plt.xticks(fontsize=15)
plt.yticks(fontsize=15)

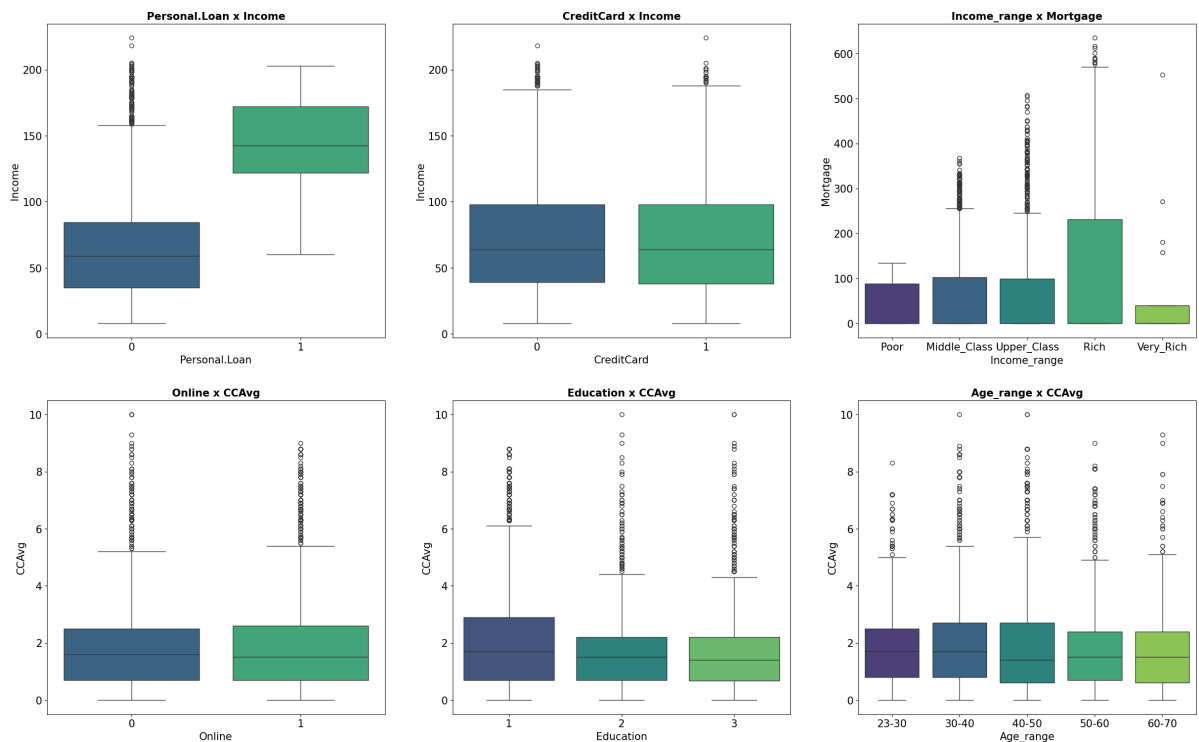
plt.subplot(3,3,4)
sns.boxplot(x='Online', y='CCAvg', data=df, palette='viridis')
plt.xlabel('Online',fontsize=15)
plt.ylabel('CCAvg',fontsize=15)
plt.title('Online x CCAvg',fontsize=15, fontweight='bold')
plt.xticks(fontsize=15)
plt.yticks(fontsize=15)

plt.subplot(3,3,5)
sns.boxplot(x='Education', y='CCAvg', data=df, palette='viridis')
plt.xlabel('Education',fontsize=15)
plt.ylabel('CCAvg',fontsize=15)
plt.title('Education x CCAvg',fontsize=15, fontweight='bold')
plt.xticks(fontsize=15)
plt.yticks(fontsize=15)

plt.subplot(3,3,6)
```

```
sns.boxplot(x='Age_r', y='CCAvg', data=df, palette='viridis')
plt.xlabel('Age_range',fontsize=15)
plt.ylabel('CCAvg',fontsize=15)
plt.title('Age_range x CCAvg',fontsize=15, fontweight='bold')
plt.xticks(fontsize=15)
plt.yticks(fontsize=15)

plt.show()
```



We can identify that income, CCAvg and Mortgage are 3 features containing outliers

Function to identify outliers

```
In [ ]: def outlier(df):

    Q1=df.quantile(0.25)

    Q3=df.quantile(0.75)

    IQR=Q3-Q1

    out = df[((df<(Q1-1.5*IQR)) | (df>(Q3+1.5*IQR)))]

    return out
```

```
In [ ]: skewed=['Income','CCAvg','Mortgage']
for col in skewed:
    outliers=outlier(df[col])
    print("Number of outliers in",col,":", str(len(outliers)),"It's Percentage is
    print("\n")
```

Number of outliers in Income : 96 ,It's Percentage is : 1.925005013033888 %

Number of outliers in CCAvg : 301 ,It's Percentage is : 6.035692801283337 %

Number of outliers in Mortgage : 291 ,It's Percentage is : 5.835171445758974 %

Outlier numbers are relatively low, yet they could better. In addition their line graphs and histograms are skewed. We found a solution to the problems by:

Find a suitable transformation for the skewed features

```
In [ ]: skewed = ['Income', 'CCAvg', 'Mortgage']

for col in skewed:
    fig, axes = plt.subplots(1, 4, figsize=(20, 5))

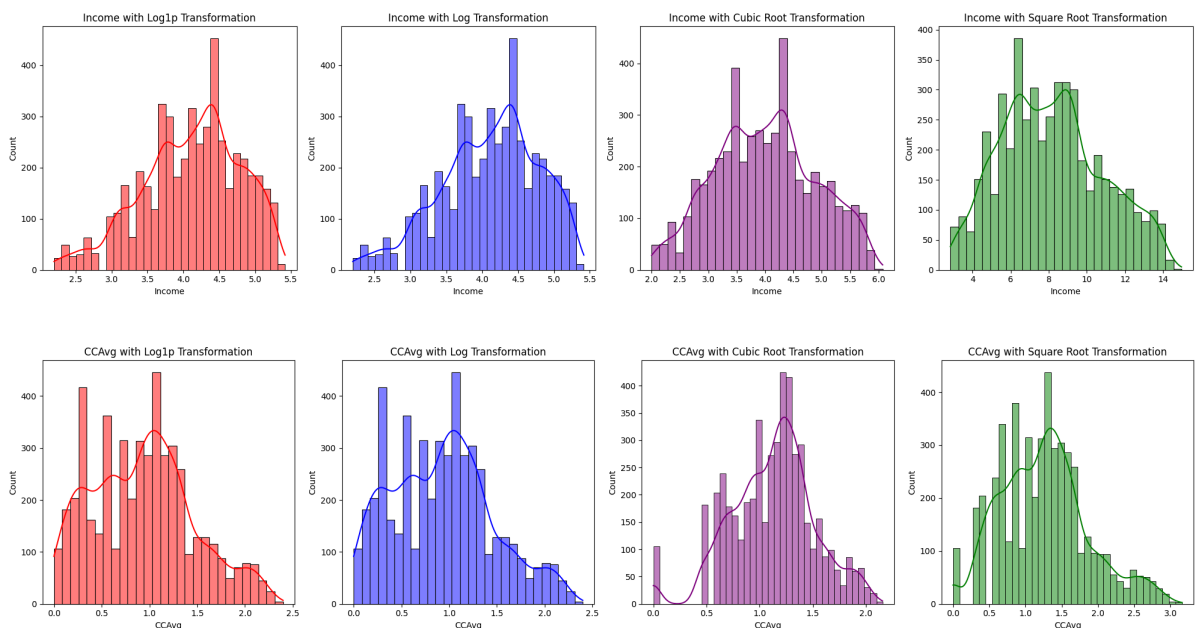
    # Log1p Transformation
    sns.histplot(np.log1p(df[col]), color='red', ax=axes[0], kde=True)
    axes[0].set_title(f'{col} with Log1p Transformation')


    # Log Transformation
    sns.histplot(np.log(df[col] + 1), color='blue', ax=axes[1], kde=True)
    axes[1].set_title(f'{col} with Log Transformation')

    # Cubic Root Transformation
    sns.histplot(np.cbrt(df[col]), color='purple', ax=axes[2], kde=True)
    axes[2].set_title(f'{col} with Cubic Root Transformation')

    # Square Root Transformation
    sns.histplot(np.sqrt(df[col]), color='green', ax=axes[3], kde=True)
    axes[3].set_title(f'{col} with Square Root Transformation')

plt.tight_layout()
plt.show()
print("\n")
```





Therefore, the best suitable transformation for:

- Income: Cubic root
- CCAvg: Cubic root
- Mortgage: Square root

```
In [ ]: df['Income'] = np.cbrt(df['Income'])
df['CCAvg'] = np.cbrt(df['CCAvg'])
df['Mortgage'] = np.sqrt(df['Mortgage'])
```

Test for outliers after the transformation and plot the histograms

```
In [ ]: skewed = ['Income', 'CCAvg', 'Mortgage']
for col in skewed:
    outliers=outlier(df[col])
    print("Number of outliers in",col,":", str(len(outliers)),",It's Percentage is ")
    print("\n")
```

Number of outliers in Income : 0 ,It's Percentage is : 0.0 %

Number of outliers in CCAvg : 109 ,It's Percentage is : 2.1856827752155605 %

Number of outliers in Mortgage : 1 ,It's Percentage is : 0.020052135552436335 %

Outliers are significantly reduced after the transformation

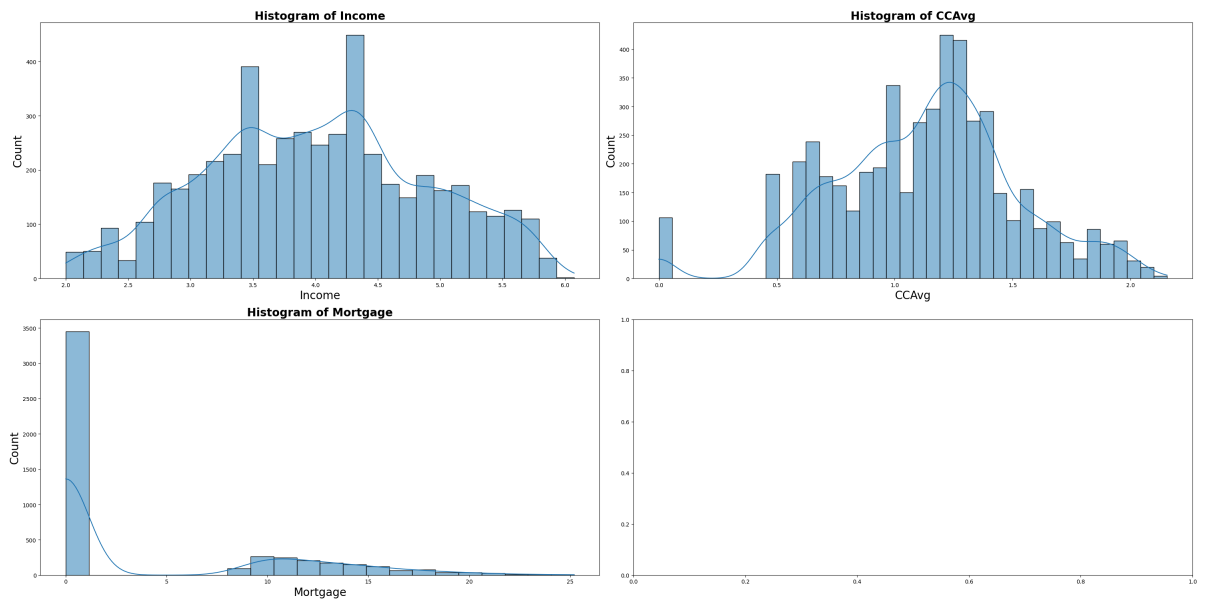
```
In [ ]: dist_columns = ['Income', 'CCAvg', 'Mortgage']

fig, axes = plt.subplots(nrows=2, ncols=2, figsize=(30, 15))

axes = axes.flatten()

for i, col in enumerate(dist_columns):
    sns.histplot(df[col], kde=True, ax=axes[i])
    axes[i].set_title(f'Histogram of {col}', fontsize=20, fontweight='bold')
    axes[i].set_xlabel(col, fontsize=20)
    axes[i].set_ylabel('Count', fontsize=20)

plt.tight_layout()
plt.show()
```



Finally, here's a summary of our continuous features

```
In [ ]: df.describe()
```

```
Out[ ]:
```

	Age	Income	CCAvg	Mortgage
<b>count</b>	4987.000000	4987.000000	4987.000000	4987.000000
<b>mean</b>	45.347704	4.006779	1.127270	4.046775
<b>std</b>	11.460838	0.887801	0.392599	6.346477
<b>min</b>	23.000000	2.000000	0.000000	0.000000
<b>25%</b>	35.000000	3.391211	0.887904	0.000000
<b>50%</b>	45.000000	4.000000	1.144714	0.000000
<b>75%</b>	55.000000	4.610436	1.375069	10.049876
<b>max</b>	67.000000	6.073178	2.154435	25.199206

## Pilot Study (Phase 3)

```
In [ ]: df1 = df.copy()

df1['Income_r'] = pd.factorize(df1['Income_r'])[0] + 1
df1['Age_r'] = pd.factorize(df1['Age_r'])[0] + 1

df1['Income_r'] = df1['Income_r'].astype(int)
df1['Age_r'] = df1['Age_r'].astype(int)

print(df1.dtypes)

df1
```

```

Age                int64
Income             float64
Family             category
CCAvg              float64
Education           category
Mortgage           float64
Personal.Loan      category
Securities.Account category
CD.Account         category
Online            category
CreditCard        category
Age_r              int64
Income_r           int64
dtype: object

```

```

Out[ ]:
   Age  Income  Family  CCAvg  Education  Mortgage  Personal.Loan  Securities.Account  C
0    25  3.659306      4  1.169607         1    0.000000           0             1
1    45  3.239612      3  1.144714         1    0.000000           0             1
2    39  2.223980      1  1.000000         1    0.000000           0             0
3    35  4.641589      1  1.392477         2    0.000000           0             0
4    35  3.556893      4  1.000000         2    0.000000           0             0
...   ...      ...      ...      ...      ...      ...      ...             ...
4995  29  3.419952      1  1.238562         3    0.000000           0             0
4996  30  2.466212      4  0.736806         1    9.219544           0             0
4997  63  2.884499      2  0.669433         3    0.000000           0             0
4998  65  3.659306      3  0.793701         2    0.000000           0             0
4999  28  4.362071      3  0.928318         1    0.000000           0             0

```

4987 rows × 13 columns



Standardize our variable

```

In [ ]: from sklearn.preprocessing import StandardScaler

standard_scaler = StandardScaler()

df_scaled=df1.copy()
columns = ['Age', 'Income', 'Family', 'CCAvg', 'Education', 'Mortgage', 'Personal.Loan', 'Securities.Account', 'CD.Account', 'Online', 'CreditCard']
for col in columns:
    df_scaled[col] = standard_scaler.fit_transform(np.array(df_scaled[col]).reshape(-1,))

df_scaled.head()

```

	Age	Income	Family	CCAvg	Education	Mortgage	Personal.Loan	Securities.Account
0	-1.775590	-0.391426	1.397399	0.107848	-1.047290	-0.637705	0	2.9246
1	-0.030341	-0.864208	0.525860	0.044436	-1.047290	-0.637705	0	2.9246
2	-0.553916	-2.008309	-1.217219	-0.324207	-1.047290	-0.637705	0	-0.3419
3	-0.902966	0.715108	-1.217219	0.675583	0.143778	-0.637705	0	-0.3419
4	-0.902966	-0.506793	1.397399	-0.324207	0.143778	-0.637705	0	-0.3419

Checking if there are any missing values

```
In [ ]: df_scaled.isnull().sum()
```

```
Out[ ]: Age                0
Income                0
Family               0
CCAvg                0
Education            0
Mortgage             0
Personal.Loan        0
Securities.Account   0
CD.Account           0
Online               0
CreditCard           0
Age_r                0
Income_r             0
dtype: int64
```

Assigning our target and decision variables

```
In [ ]: Y = df_scaled['Personal.Loan']
X = df_scaled.drop(['Personal.Loan'],axis=1)
```

```
In [ ]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, Y, test_size=0.4, random_state=42)
```

```
In [ ]: print (" Number of columns in our Features : ", X.shape[1])

Number of columns in our Features : 12
```

Solving the Class imbalance problem

```
In [ ]: from imblearn.over_sampling import SMOTE

smote = SMOTE(random_state=42)
X_train_upsampled, y_train_upsampled = smote.fit_resample(X_train, y_train)
```

```
In [ ]: print("Before UpSampling, counts of Personal loan = '0': {}".format(sum(y_train==0)))
print("Before UpSampling, counts of Personal loan = '1': {}".format(sum(y_train==1)))

print("After UpSampling, counts of Personal loan = '0': {}".format(sum(y_train_upsampled==0)))
print("After UpSampling, counts of Personal loan = '1': {}".format(sum(y_train_upsampled==1)))
```



Before UpSampling, counts of Personal loan = '0': 2699  
Before UpSampling, counts of Personal loan = '1': 293

After UpSampling, counts of Personal loan = '0': 2699  
After UpSampling, counts of Personal loan = '1': 2699

Initialize a Data Frame to store the Accuracy, Precision, Recall, and F1 score for all our upcoming model

```
In [ ]: EVAL_SCORE = pd.DataFrame(columns=['Model', 'Accuracy', 'Precision', 'Recall', 'F1 Score'])  
EVAL_SCORE
```

```
Out[ ]: 

| Model | Accuracy | Precision | Recall | F1 Score |
|-------|----------|-----------|--------|----------|
|-------|----------|-----------|--------|----------|


```

Decision Tree Model

```
In [ ]: from sklearn.tree import DecisionTreeClassifier  
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix  
from sklearn.model_selection import cross_val_score  
import plotly.graph_objects as go
```

```
In [ ]: max_depth_values = range(1, 50)  
  
train_scores = []  
test_scores = []  
  
for depth in max_depth_values:  
    clf = DecisionTreeClassifier(max_depth=depth, random_state=42)  
    clf.fit(X_train_upsampled, y_train_upsampled)  
  
    y_train_pred = clf.predict(X_train_upsampled)  
    train_scores.append(accuracy_score(y_train_upsampled, y_train_pred))  
  
    y_test_pred = clf.predict(X_test)  
    test_scores.append(accuracy_score(y_test, y_test_pred))  
  
fig = go.Figure()  
  
fig.add_trace(go.Scatter(x=list(max_depth_values), y=train_scores, mode='lines', name='Train Accuracy'))  
fig.add_trace(go.Scatter(x=list(max_depth_values), y=test_scores, mode='lines', name='Test Accuracy'))  
  
fig.update_layout(  
    title='Max Depth vs. Accuracy',  
    xaxis=dict(title='Max Depth'),  
    yaxis=dict(title='Accuracy'),  
    legend=dict(x=0.7, y=0.9),  
)  
  
fig.show()
```

From this graph we could deduce that after a depth=3 the graph starts to flatten out; therefore, we'd build our decision tree model using a maximum depth=3. That would allow us to overcome overfitting problems

```
In [ ]: min_samples_split_values = range(2, 30)

train_scores = []
test_scores = []

for split in min_samples_split_values:
    clf = DecisionTreeClassifier(min_samples_split=split, random_state=42)
    clf.fit(X_train_upsampled, y_train_upsampled)

    y_train_pred = clf.predict(X_train_upsampled)
    train_scores.append(accuracy_score(y_train_upsampled, y_train_pred))

    y_test_pred = clf.predict(X_test)
    test_scores.append(accuracy_score(y_test, y_test_pred))

fig = go.Figure()

fig.add_trace(go.Scatter(x=list(min_samples_split_values), y=train_scores, mode='li
fig.add_trace(go.Scatter(x=list(min_samples_split_values), y=test_scores, mode='li

fig.update_layout(
    title='Min Samples Split vs. Accuracy',
    xaxis=dict(title='Min Samples Split'),
    yaxis=dict(title='Accuracy'),
```

```

        legend=dict(x=0.7, y=0.9),
    )

fig.show()

```

From the graph above we could conclude that the optimal minimum sample split is =3 where it provides the maximum accuracy

```
In [ ]: Decision_Tree = DecisionTreeClassifier(max_depth=3,criterion='entropy',random_state=42)
```

```
In [ ]: Decision_Tree.fit(X_train_upsampled, y_train_upsampled)
```

```
Out[ ]: ▼ DecisionTreeClassifier
DecisionTreeClassifier(criterion='entropy', max_depth=3, random_state=42)
```

```
In [ ]: y_pred_train = Decision_Tree.predict(X_train_upsampled)
y_pred_test = Decision_Tree.predict(X_test)
```

```
In [ ]: train_accuracy = accuracy_score(y_train_upsampled, y_pred_train)
print(" Decision Tree Training Accuracy :",round(train_accuracy,2)*100)

test_accuracy = accuracy_score(y_test, y_pred_test)
print(" Decision Tree Testing Accuracy :",round(test_accuracy,2)*100)
```

```

Decision Tree Training Accuracy : 98.0
Decision Tree Testing Accuracy : 96.0

```

Using cross validation on our decision tree and testing the accuracy

```
In [ ]: cv_scores_train = cross_val_score(Decision_Tree, X_train_upsampled, y_train_upsampled)
print("Cross-Validation Scores on Training Data: ", cv_scores_train)
print(" Mean Accuracy from Cross-Validation : ", cv_scores_train.mean())
```

Cross-Validation Scores on Training Data: [0.96666667 0.96944444 0.97685185 0.97034291 0.97126969]  
Mean Accuracy from Cross-Validation : 0.970915113445234

```
In [ ]: conf_matrix = confusion_matrix(y_test, y_pred_test)

# Add labels for better understanding
tn, fp, fn, tp = conf_matrix.ravel()
display( pd.DataFrame(conf_matrix, columns=['Predicted Negative', 'Predicted Positive'])
```

	Predicted Negative	Predicted Positive
Actual Negative	1741	67
Actual Positive	3	184

```
In [ ]: print("Classification Report : \n" ,classification_report(y_test, y_pred_test))
```

```
Classification Report :
              precision    recall  f1-score   support

     0           1.00       0.96       0.98       1808
     1           0.73       0.98       0.84        187

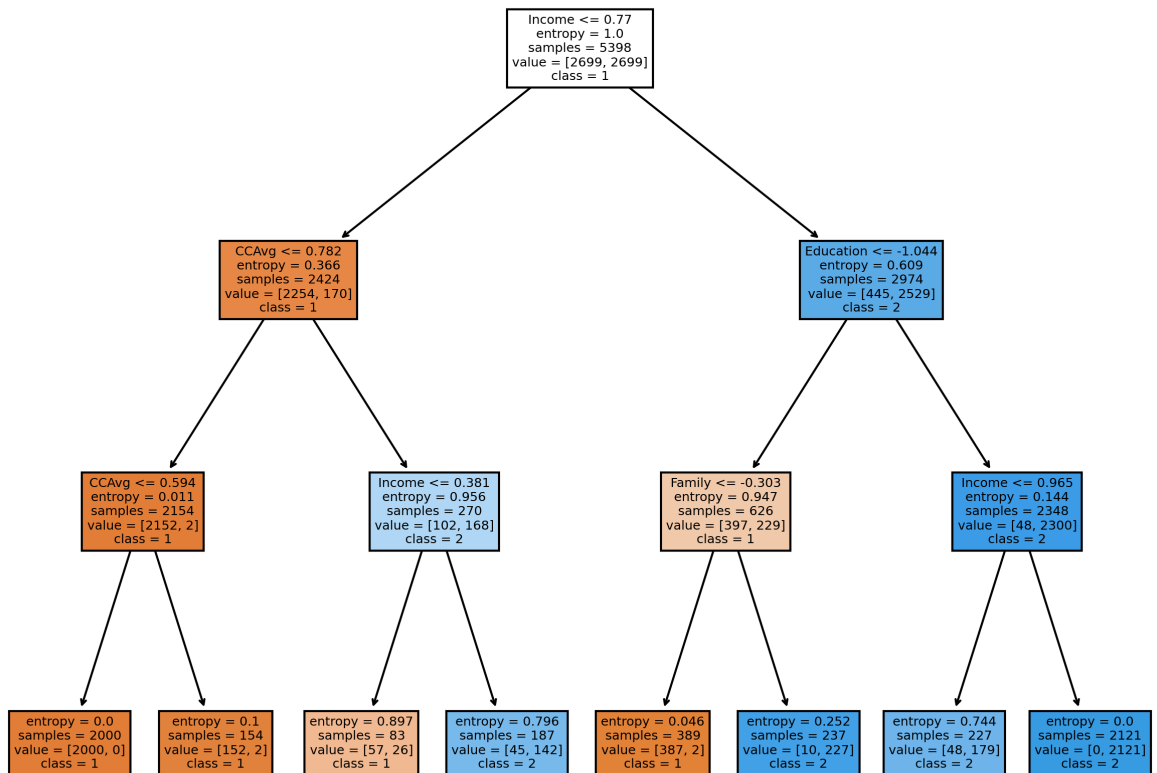
 accuracy              0.96       1995
 macro avg           0.87       0.97       0.91       1995
 weighted avg        0.97       0.96       0.97       1995
```

```
In [ ]: from sklearn.tree import plot_tree
import matplotlib.pyplot as plt

cn=["1","2"]
fn=['Age',      'Income',      'Family',      'CCAvg',      'Education',      'Mc

fig, axes = plt.subplots(nrows = 1,ncols = 1,figsize = (10,8), dpi=300)
plot_tree(Decision_Tree, filled=True,feature_names = fn,class_names=cn)
plt.title("Decision Tree Model")
plt.show()
```

## Decision Tree Model



```

In [ ]: accuracy = accuracy_score(y_test, y_pred_test)
precision = precision_score(y_test, y_pred_test, average='macro')
recall = recall_score(y_test, y_pred_test, average='macro')
f1_score = metrics.f1_score(y_test, y_pred_test, average='macro')

EVAL_SCORE = EVAL_SCORE.append({'Model': 'Decision Tree', 'Accuracy': accuracy, 'Precision': precision, 'Recall': recall, 'F1 Score': f1_score})

```

```

Out[ ]:

```

	Model	Accuracy	Precision	Recall	F1 Score
0	Decision Tree	0.964912	0.865674	0.97345	0.910238

## Random Forest Model

```

In [ ]: from sklearn.ensemble import RandomForestClassifier

clf = RandomForestClassifier(max_depth=2, random_state=0)

clf.fit(X_train, y_train)

```

```

Out[ ]:
RandomForestClassifier
RandomForestClassifier(max_depth=2, random_state=0)

```

```

In [ ]: y_pred_train_rf = clf.predict(X_train)
y_pred_test_rf = clf.predict(X_test)

```

```

In [ ]: train_accuracy = accuracy_score(y_train, y_pred_train_rf)
print("Random Forest Training Accuracy :", round(train_accuracy, 3) * 100)

```

```
test_accuracy = accuracy_score(y_test, y_pred_test_rf)
print("Random Forest Testing Accuracy : ",round(test_accuracy,3)*100)
```

Random Forest Training Accuracy : 91.7  
Random Forest Testing Accuracy : 91.9

```
In [ ]: conf_matrix = confusion_matrix(y_test, y_pred_test_rf)

# Add Labels for better understanding
tn, fp, fn, tp = conf_matrix.ravel()
display( pd.DataFrame(conf_matrix, columns=['Predicted Negative', 'Predicted Positive'])
```

	Predicted Negative	Predicted Positive
Actual Negative	1808	0
Actual Positive	162	25

```
In [ ]: print("Classification Report : \n" ,classification_report(y_test, y_pred_test_rf))
```

```
Classification Report :
              precision    recall  f1-score   support

      0       0.92      1.00      0.96      1808
      1       1.00      0.13      0.24       187

 accuracy          0.96
 macro avg         0.96
weighted avg         0.93
```

```
In [ ]: accuracy = accuracy_score(y_test, y_pred_test_rf)
precision = precision_score(y_test, y_pred_test_rf, average='macro')
recall = recall_score(y_test, y_pred_test_rf, average='macro')
f1_score= metrics.f1_score(y_test, y_pred_test_rf, average='macro')

EVAL_SCORE = EVAL_SCORE.append({'Model': 'Random Forest', 'Accuracy':accuracy, 'Precision':precision, 'Recall':recall, 'F1 Score':f1_score})
EVAL_SCORE
```

```
Out[ ]:
```

	Model	Accuracy	Precision	Recall	F1 Score
0	Decision Tree	0.964912	0.865674	0.973450	0.910238
1	Random Forest	0.918797	0.958883	0.566845	0.596485

Logistic Regression Model

```
In [ ]: from sklearn.linear_model import LogisticRegression
logistic_model=LogisticRegression()
logistic_model.fit(X_train,y_train)
```

```
Out[ ]: ▾ LogisticRegression
LogisticRegression()
```

```
In [ ]: y_pred_log=logistic_model.predict(X_test)

y_pred_train_log=logistic_model.predict(X_train)
```

```
In [ ]: conf_matrix = confusion_matrix(y_test, y_pred_log)
```

```
# Add Labels for better understanding
tn, fp, fn, tp = conf_matrix.ravel()
display(pd.DataFrame(conf_matrix, columns=['Predicted Negative', 'Predicted Positive']))
```

	Predicted Negative	Predicted Positive
Actual Negative	1783	25
Actual Positive	56	131

```
In [ ]: print("Classification Report : \n", classification_report(y_test, y_pred_log))
```

```
Classification Report :
              precision    recall  f1-score   support

     0       0.97       0.99       0.98       1808
     1       0.84       0.70       0.76        187

 accuracy          0.96          1995
 macro avg         0.90          1995
 weighted avg      0.96          1995
```

```
In [ ]: from sklearn.metrics import accuracy_score, confusion_matrix, classification_report

print(f"Logistic Regression training Accuracy score: {accuracy_score(y_train,y_pred_log)}")
print(f"Logistic Regression testing Accuracy score: {accuracy_score(y_test,y_pred_log)}")
```

```
Logistic Regression training Accuracy score: 95.9
Logistic Regression testing Accuracy score: 95.9
```

Plot the ROC curve to better identify the threshold by which we would split our data

```
In [ ]: from sklearn.metrics import roc_curve, auc

train_probs = logistic_model.predict_proba(X_train)[: , 1]
test_probs = logistic_model.predict_proba(X_test)[: , 1]

fpr_train, tpr_train, thresholds_train = roc_curve(y_train, train_probs)
roc_auc_train = auc(fpr_train, tpr_train)


fpr_test, tpr_test, thresholds_test = roc_curve(y_test, test_probs)
roc_auc_test = auc(fpr_test, tpr_test)

plt.figure(figsize=(8, 6))
plt.plot(fpr_train, tpr_train, label=f'Train ROC curve (AUC = {roc_auc_train:.2f})')
# Plot ROC curve for test set
plt.plot(fpr_test, tpr_test, label=f'Test ROC curve (AUC = {roc_auc_test:.2f})')

for i, threshold in enumerate(thresholds_test):
    if i % 50 == 0:
        plt.annotate(f'{threshold:.2f}', (fpr_test[i], tpr_test[i]), textcoords="offsetpoints",
                    dx=10, dy=-10)

plt.plot([0, 1], [0, 1], 'r--', label='Random Classifier')

plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve')
plt.legend()
plt.show()
```



Therefore, the best threshold would be at 0.04

```
In [ ]: threshold=0.04
test_prob=logistic_model.predict_proba(X_test)
y_pred_test_log = (test_prob[:, 1] >= threshold).astype(int)

In [ ]: print(f"Logistic Regression testing Accuracy score: {accuracy_score(y_test,y_pred_test_log)}")
Logistic Regression testing Accuracy score: 85.8

In [ ]: conf_matrix = confusion_matrix(y_test, y_pred_test_log)

# Add labels for better understanding
tn, fp, fn, tp = conf_matrix.ravel()
display(pd.DataFrame(conf_matrix, columns=['Predicted Negative', 'Predicted Positive'], index=['Actual Negative', 'Actual Positive']))
```

	Predicted Negative	Predicted Positive
Actual Negative	1535	273
Actual Positive	10	177

```
In [ ]: print("Classification Report : \n", classification_report(y_test, y_pred_test_log))
```



```

Classification Report :
              precision    recall  f1-score   support

     0       0.99      0.85      0.92      1808
     1       0.39      0.95      0.56       187

 accuracy          0.86      1995
 macro avg         0.69      1995
 weighted avg      0.94      1995

```

```

In [ ]: accuracy = accuracy_score(y_test, y_pred_test_log)
precision = precision_score(y_test, y_pred_test_log, average='macro')
recall = recall_score(y_test, y_pred_test_log, average='macro')
f1_score = metrics.f1_score(y_test, y_pred_test_log, average='macro')

EVAL_SCORE = EVAL_SCORE.append({'Model': 'Logistic Regression', 'Accuracy': accuracy,
                                'Precision': precision, 'Recall': recall, 'F1 Score': f1_score})

```

```

Out[ ]:

```

	Model	Accuracy	Precision	Recall	F1 Score
0	Decision Tree	0.964912	0.865674	0.973450	0.910238
1	Random Forest	0.918797	0.958883	0.566845	0.596485
2	Logistic Regression	0.858145	0.693430	0.897764	0.735664

#### K-NN Classifier Model

```

In [ ]: from sklearn.neighbors import KNeighborsClassifier
from sklearn.model_selection import cross_val_predict, KFold

knn = KNeighborsClassifier(n_neighbors = 5)

```

```

In [ ]: knn.fit(X_train, y_train)

```

```

Out[ ]:
KNeighborsClassifier
KNeighborsClassifier()

```

```

In [ ]: kf = KFold(n_splits=10, random_state=5, shuffle=True)
y_pred_knn = cross_val_predict(knn, X, Y, cv=kf)

```

```

In [ ]: knn.score(X_test, y_test)

```

```

Out[ ]: 0.9588972431077695

```

```

In [ ]: k_values = range(1, 50)

#intialize 2 variables
best_k = 0
best_accuracy = 0

#iterate over the values of k and calculate the cross-validation for each, then cal
#accuracy and compare

for k in k_values:
    knn = KNeighborsClassifier(n_neighbors=k)
    kf = KFold(n_splits=10, random_state=k, shuffle=True)
    y_pred = cross_val_predict(knn, X_test, y_test , cv=kf)

```

```

accuracy = accuracy_score(y_test, y_pred)
if accuracy > best_accuracy:
    best_accuracy = accuracy
    best_k = k
#use cross-validation for the best k provided from the loop
final_knn = KNeighborsClassifier(n_neighbors=best_k)
kf = KFold(n_splits=10, random_state=best_k, shuffle=True)
y_pred_KNN = cross_val_predict(final_knn, X_test, y_test, cv=kf)

```

```

In [ ]: accuracy_final = accuracy_score(y_test, y_pred_KNN)
precision_final = precision_score(y_test, y_pred_KNN, average='macro')
recall_final = recall_score(y_test, y_pred_KNN, average='macro')
f1_score_final = metrics.f1_score(y_test, y_pred_KNN, average='macro')

print(f"KNN Classification Results (Best k={best_k}):")
print(f"Accuracy: {accuracy_final:.4f}")
print(f"Precision: {precision_final:.4f}")
print(f"Recall: {recall_final:.4f}")
print(f"F1 Score: {f1_score_final:.4f}")

```

```

KNN Classification Results (Best k=1):
Accuracy: 0.9564
Precision: 0.8843
Recall: 0.8489
F1 Score: 0.8655

```

```

In [ ]: conf_matrix = confusion_matrix(y_test, y_pred_KNN)

# Add Labels for better understanding
tn, fp, fn, tp = conf_matrix.ravel()
display(pd.DataFrame(conf_matrix, columns=['Predicted Negative', 'Predicted Positive']

```

	Predicted Negative	Predicted Positive
Actual Negative	1774	34
Actual Positive	53	134

```

In [ ]: EVAL_SCORE = EVAL_SCORE.append({'Model': 'K-NN', 'Accuracy': accuracy_final, 'Precision': precision_final, 'Recall': recall_final, 'F1 Score': f1_score_final})
EVAL_SCORE

```

```

Out[ ]:

```

	Model	Accuracy	Precision	Recall	F1 Score
0	Decision Tree	0.964912	0.865674	0.973450	0.910238
1	Random Forest	0.918797	0.958883	0.566845	0.596485
2	Logistic Regression	0.858145	0.693430	0.897764	0.735664
3	K-NN	0.956391	0.884305	0.848886	0.865498

Linear Regression Model

```

In [ ]: from sklearn.linear_model import LinearRegression

linreg = LinearRegression()
linreg.fit(X_train, y_train)

```

```

Out[ ]:
LinearRegression()

```

```
In [ ]: from sklearn.metrics import mean_squared_error
```

```
print('linear model coeff (w): {}'.format(linreg.coef_))
print('linear model intercept (b): {:.3f}'.format(linreg.intercept_))
print('R-squared score (training): {:.3f}'.format(linreg.score(X_train, y_train)))
print('R-squared score (test): {:.3f}'.format(linreg.score(X_test, y_test)))

y_pred_LR=linreg.predict(X_test)
print('MSE (test): {:.3f}'.format(mean_squared_error(y_test, y_pred_LR)))
```

```
linear model coeff (w): [ 0.00985661  0.07236611  0.0357375   0.01712037  0.072341
 87  0.00831916
 -0.02275375  0.08017948 -0.01209817 -0.02220952 -0.00439635  0.07719119]
linear model intercept (b): -0.024
R-squared score (training): 0.385
R-squared score (test): 0.399
MSE (test): 0.051
```

We won't be able to add it to our evaluation table since linear regression is assessed by MSE and  $R^2$  which are provided above rather than accuracy, precision, recall, and F1 score

#### Naive Bayes

```
In [ ]: from sklearn.naive_bayes import GaussianNB
```

```
nbclf = GaussianNB()
nbclf.fit(X_train, y_train)
```

```
Out[ ]: ▾ GaussianNB
GaussianNB()
```

```
In [ ]: print('Accuracy of GaussianNB classifier on training set: {:.2f}'
            .format(nbclf.score(X_train, y_train)))
print('Accuracy of GaussianNB classifier on test set: {:.2f}'
            .format(nbclf.score(X_test, y_test)))
```

```
Accuracy of GaussianNB classifier on training set: 0.88
Accuracy of GaussianNB classifier on test set: 0.90
```

```
In [ ]: y_pred_NB=nbclf.predict(X_test)
```

```
In [ ]: conf_matrix = confusion_matrix(y_test, y_pred_NB)

# Add labels for better understanding
tn, fp, fn, tp = conf_matrix.ravel()
display(pd.DataFrame(conf_matrix, columns=['Predicted Negative', 'Predicted Positive']
```

	Predicted Negative	Predicted Positive
Actual Negative	1633	175
Actual Positive	34	153

```
In [ ]: accuracy = accuracy_score(y_test, y_pred_NB)
precision = precision_score(y_test, y_pred_NB, average='macro')
```

```
recall = recall_score(y_test, y_pred_NB, average='macro')
f1_score= metrics.f1_score(y_test, y_pred_NB, average='macro')

EVAL_SCORE = EVAL_SCORE.append({'Model': 'Naive Bayes', 'Accuracy':accuracy, 'Precision':precision, 'Recall':recall, 'F1 Score':f1_score})
EVAL_SCORE
```

```
Out[ ]:
```

	Model	Accuracy	Precision	Recall	F1 Score
0	Decision Tree	0.964912	0.865674	0.973450	0.910238
1	Random Forest	0.918797	0.958883	0.566845	0.596485
2	Logistic Regression	0.858145	0.693430	0.897764	0.735664
3	K-NN	0.956391	0.884305	0.848886	0.865498
4	Naive Bayes	0.895238	0.723034	0.860695	0.767015

## Assessing our Models

```
In [ ]: EVAL_SCORE
```

```
Out[ ]:
```

	Model	Accuracy	Precision	Recall	F1 Score
0	Decision Tree	0.964912	0.865674	0.973450	0.910238
1	Random Forest	0.918797	0.958883	0.566845	0.596485
2	Logistic Regression	0.858145	0.693430	0.897764	0.735664
3	K-NN	0.956391	0.884305	0.848886	0.865498
4	Naive Bayes	0.895238	0.723034	0.860695	0.767015

Sorting our evaluation table descendingly according to the F1 score

```
In [ ]: EVAL_SCORE.sort_values(by='F1 Score', inplace = True, ascending= False)
EVAL_SCORE
```

```
Out[ ]:
```

	Model	Accuracy	Precision	Recall	F1 Score
0	Decision Tree	0.964912	0.865674	0.973450	0.910238
3	K-NN	0.956391	0.884305	0.848886	0.865498
4	Naive Bayes	0.895238	0.723034	0.860695	0.767015
2	Logistic Regression	0.858145	0.693430	0.897764	0.735664
1	Random Forest	0.918797	0.958883	0.566845	0.596485

```
In [ ]: plt.figure(figsize=(30,28))
plt.subplot(2,2,1)

acc = EVAL_SCORE.groupby('Model')['Accuracy'].mean()
acc.plot(kind='bar',color=sns.palettes.mpl_palette('Dark2'))
plt.xticks(fontsize=15)
plt.ylabel('Accuracy',fontsize=15)
plt.title('Model Accuracy',fontsize=20, fontweight="bold")

plt.subplot(2,2,2)
```

```

pr = EVAL_SCORE.groupby('Model')['Precision'].mean()
pr.plot(kind='bar',color=sns.palettes.mpl_palette('Dark2'))
plt.xticks(fontsize=15)
plt.ylabel('Precision',fontsize=15)
plt.title('Model Precision',fontsize=20, fontweight="bold")

plt.subplot(2,2,3)

rc = EVAL_SCORE.groupby('Model')['Recall'].mean()
rc.plot(kind='bar',color=sns.palettes.mpl_palette('Dark2'))
plt.xticks(fontsize=15)
plt.ylabel('Recall',fontsize=15)
plt.title('Model Recall',fontsize=20, fontweight="bold")

plt.subplot(2,2,4)

f1 = EVAL_SCORE.groupby('Model')['F1 Score'].mean()
f1.plot(kind='bar',color=sns.palettes.mpl_palette('Dark2'))
plt.xticks(fontsize=15)
plt.ylabel('F1 Score',fontsize=15)
plt.title('Model F1 Score',fontsize=20, fontweight="bold")

plt.subplots_adjust(wspace=0.5,hspace=0.5)
plt.show()

```

